

# **CS 5787 Assignment 2**

**Ran Ji**  
**rj369@**

# Problem 1

## Part 1 - Using Pre-Trained Deep CNN (5 points)

Model used: resnet152

```
1 normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
2
3 train_loader = torch.utils.data.DataLoader(
4     datasets.ImageFolder("./q1_train_img",
5         transforms.Compose([transforms.ToTensor(),normalize])),
6     batch_size=10)
7
8 pepper_image = Variable(next(iter(train_loader))[0])
9
10 resnet152 = models.resnet152(pretrained=True)
11 resnet152.eval()
12
13 resnet152.conv1.register_forward_hook(get_activation('early'))
14 resnet152.layer2[1].conv3.register_forward_hook(get_activation('middle'))
15 resnet152.layer4[1].conv3.register_forward_hook(get_activation('end'))
16
17 output = resnet152(pepper_image)
18
19 prob = sum(np.exp(output[0].detach().numpy()))
20 index_file = json.load(open("./imagenet_class_index.txt"))
21
22 index_name = [0]*len(index_file)
23 for i in range(len(index_file)):
24     index_name[i] = index_file[str(i)][1]
```

```
1 output2 = output[0].sort()[1][-3:]
2 output2
```

tensor([831, 943, 945])

```
1 for index in reversed(output2):
2     print(index_name[index], end = "")
3     print(" probability is ", end = "")
4     print(np.exp(output[0][index].detach().numpy())/prob)
```

bell\_pepper probability is 0.9002557140463503  
cucumber probability is 0.08118853944603413  
studio\_couch probability is 0.005074153097658631


## Part 2 - Visualizing Feature Maps (5 points)

bell\_pepper probability is 0.9002557140463503  
cucumber probability is 0.08118853944603413  
studio\_couch probability is 0.005074153097658631

```
1 print(activation['early'].shape)
2 print(activation['middle'].shape)
3 print(activation['end'].shape)
```

```
torch.Size([1, 64, 112, 112])
torch.Size([1, 512, 28, 28])
torch.Size([1, 2048, 7, 7])
```

```
1 def print_feature(layer):
2     layer = layer - layer.min()
3     layer = layer / layer.max()
4     fig = plt.figure(figsize=(12,8))
5     for i in range(5):
6         image = layer[i]
7         fig.add_subplot(1, 5, i+1)
8         plt.imshow(image)
9     plt.show()
10
11 print_feature(activation['early'][0])
12 print_feature(activation['middle'][0])
13 print_feature(activation['end'][0])
```



## Problem 2

### Extract features:

```
1 import numpy as np
2 import json
3 import torch
4 import torch.nn as nn
5 import torch.nn.functional as F
6 import random
7 import torch.optim as optim
8
9 import torchvision.models as models
10 import torchvision.transforms as transforms
11 import torchvision.datasets as datasets
12 from torch.autograd import Variable
13 import matplotlib.pyplot as plt
14 import matplotlib.image as mpimg
15 import matplotlib
16
17 from PIL import Image
18 from random import shuffle
```

```
1 content = open("annotations/trainval.txt","r").readlines()
2 PERCENT = 20
3
4 train_data = {}
5 for line in content:
6     if line.strip() == "":
7         continue
8     if random.randint(1,PERCENT) % PERCENT != 0: # remove
9         continue
10    name = line.split()[0].strip()
11    value = int(line.split()[1].strip())
12    train_data[name] = value
13
14
15 content = open("annotations/test.txt","r").readlines()
16
17 test_data = {}
18 for line in content:
19     if line.strip() == "":
20         continue
21     # if random.randint(1,PERCENT) % PERCENT != 0: # remove
22     #     continue
23     name = line.split()[0].strip()
24     value = int(line.split()[1].strip())
25     test_data[name] = value
26
```

```
1 FEATURE_NUM = 2048
2 CLASS_NUM = 37
3
4 train_features = np.zeros((len(train_data), FEATURE_NUM))
5 train_label = []
6 train_items = list(train_data.items())
7
8 test_features = np.zeros((len(test_data), FEATURE_NUM))
9 test_label = []
10 test_items = list(test_data.items())
11
12 model = models.resnet152(pretrained=True)
13 extra_feature_model = nn.Sequential(*list(model.children()))[:-1]
```

```

1 def extract_features(is_train):
2     global model
3
4     batch_size = 20
5     items = None
6     features = None
7     labels = None
8
9     if (is_train):
10        items = train_items
11        features = train_features
12        labels = train_label
13    else:
14        items = test_items
15        features = test_features
16        labels = test_label
17
18    trans = transforms.Compose([
19        transforms.Resize(256),
20        transforms.CenterCrop(224),
21        transforms.ToTensor(),
22        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
23    ])
24
25    for i in range(0, len(items), batch_size):
26        current_items = items[i:i+batch_size]
27        batch_input = None
28        for name, value in current_items:
29            path = "images/"+name+".jpg"
30            image = Image.open(path).convert('RGB')
31            labels.append(value)
32            tensor = trans(image).unsqueeze(0)
33            if batch_input is None:
34                batch_input = tensor
35            else:
36                batch_input = torch.cat((batch_input, tensor), 0)
37
38        input_var = torch.autograd.Variable(batch_input)
39        output = extra_feature_model(input_var)
40        features[i:i+batch_size] = output.data.numpy().reshape((-1, FEATURE_NUM))
41        if (i%100 == 0):
42            print(i)

```

```

1 extract_features(True)
2 extract_features(False)

```

```

1 print(train_features.shape)
2 print(test_features.shape)

```

```

torch.Size([3680, 2048])
torch.Size([3669, 2048])

```

## Train linear classier and classify test features:

```

1 # L2 normalize
2 train_features = F.normalize(Variable(torch.from_numpy(train_features)), p=2, dim=1)
3 test_features = F.normalize(Variable(torch.from_numpy(test_features)), p=2, dim=1)
4
5 train_label = torch.from_numpy(np.asarray(train_label)).long()
6 test_label = np.asarray(test_label)

```

```

1 linear_classifier = LinearDiscriminantAnalysis()
2 linear_classifier.fit(train_features, train_label)
3 print(linear_classifier.score(test_features, test_label))
4 # linear_classifier.score(train_features, train_label)
5 final_result = linear_classifier.predict(test_features)

```

```
0.17879531207413465
```



## Mean-per-class accuracy:

```
1 mean_per_class = {}

1 TEST_NUM = len(final_result)
2 for i in range(TEST_NUM):
3     class_num = final_result[i]
4     result = 0
5     if test_label[i] == final_result[i]:
6         result = 1
7
8     if class_num in mean_per_class:
9         mean_per_class[class_num].append(result)
10    else:
11        mean_per_class[class_num] = [result]

1 for key in range(38):
2     if key in mean_per_class:
3         print("class " + str(key) + " accuracy is " + str(sum(mean_per_class[key])/len(mean_per_class[key])))
```

```
class 1 accuracy is 0.17
class 2 accuracy is 0.10810810810810811
class 3 accuracy is 0.07142857142857142
class 4 accuracy is 0.20930232558139536
class 5 accuracy is 0.18181818181818182
class 6 accuracy is 0.15217391304347827
class 7 accuracy is 0.1743119266055046
class 8 accuracy is 0.0967741935483871
class 9 accuracy is 0.17557251908396945
class 10 accuracy is 0.16161616161616163
class 11 accuracy is 0.10204081632653061
class 12 accuracy is 0.23958333333333334
class 13 accuracy is 0.16666666666666666
class 14 accuracy is 0.12745098039215685
class 15 accuracy is 0.21621621621621623
class 16 accuracy is 0.1724137931034483
class 17 accuracy is 0.18691588785046728
class 18 accuracy is 0.2807017543859649
class 19 accuracy is 0.2553191489361702
class 20 accuracy is 0.21904761904761905
class 21 accuracy is 0.15384615384615385
class 22 accuracy is 0.18518518518518517
class 23 accuracy is 0.22641509433962265
class 24 accuracy is 0.18181818181818182
class 25 accuracy is 0.175
class 26 accuracy is 0.25688073394495414
class 27 accuracy is 0.11206896551724138
class 28 accuracy is 0.09523809523809523
class 29 accuracy is 0.22522522522522523
class 30 accuracy is 0.25
class 31 accuracy is 0.2169811320754717
class 32 accuracy is 0.22448979591836735
class 33 accuracy is 0.15555555555555556
class 34 accuracy is 0.18518518518518517
class 35 accuracy is 0.09090909090909091
class 36 accuracy is 0.2619047619047619
class 37 accuracy is 0.15555555555555556
```

## Classifier used:

I used the LinearDiscriminantAnalysis classifier.


## Problem 3

### Part 1

#### IMAGE SHOWING THE FILTERS



```
: 1 def print_filter():
2     print(model.layer1.weight.detach().shape)
3     layer = model.layer1.weight.detach().permute(0,2,3,1)
4     layer = layer - layer.min()
5     layer = layer / layer.max()
6     fig=plt.figure(figsize=(15,5))
7     for i in range(64):
8         f = layer[i]
9         fig.add_subplot(4,16, i+1)
10        plt.imshow(f)
11    plt.show()
12
13 print_filter()
```

torch.Size([64, 3, 11, 11])



## TRAINING LOSS AS FUNCTION OF EPOCHS

```
1 plt.subplots(figsize=(7,5))
2 plt.plot(range(epoch_num), train_acc, label="train set")
3 plt.ylabel('accuracy')
4 plt.xlabel('epoch')
5 plt.legend()
6 plt.show()
7
8 plt.subplots(figsize=(7,5))
9 plt.plot(range(epoch_num), train_error, label="train set")
10 plt.ylabel('cross-entropy loss')
11 plt.xlabel('epoch')
12 plt.legend()
13 plt.show()
```





## TEST DATA ACCURACY

```
1 # run on test data
2
3 loss = nn.CrossEntropyLoss()
4 optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
5
6 test_correct_num = 0
7 test_total_num = 0
8 test_error = 0
9 for i, data in enumerate(test_loader, 0):
10     # if random.randint(1,1000) % 1000 != 0: # remove
11     #     continue
12     if (i%1000 == 0):
13         print(i)
14
15     inputs, labels = data
16
17     optimizer.zero_grad()
18     outputs = model(inputs)
19     current_loss = loss(outputs, labels)
20     _, result = torch.max(outputs.data, 1)
21
22     test_correct_num = test_correct_num + (result == labels).sum().item()
23     test_total_num = test_total_num + labels.size(0)
24     test_error = test_error + current_loss.data.item()
25
26 print("Test data acc is " + str(test_correct_num/test_total_num) + ", error is " + str(test_error/test_total_num))
0
Test data acc is 0.6791, error is 0.0927662309229374
```

---

## WEIGHT INITIALIZATION INFORMATION

The weight is initialized randomly.

## DESCRIBE HYPER-PARAMETERS

learning rate=0.01

momentum=0.9

epoch\_num = 10

batch\_size = 10

num\_workers = 2

## Full Code

```
1 import numpy as np
2 import random
3
4 import torch
5 import torch.optim as optim
6 import torch.nn as nn
7 import torch.nn.functional as F
8
9 import torchvision
10 import torchvision.models as models
11 import torchvision.transforms as transforms
12 import torchvision.datasets as datasets
13 from torch.autograd import Variable
14 import matplotlib.pyplot as plt

```

---

```
1 transform = transforms.Compose([
2     transforms.ToTensor(),
3     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))
4 ])
5
6 train_data = torchvision.datasets.CIFAR10(root='./cifar-10-python/train', train=True, download=True, transform=transform)
7 train_loader = torch.utils.data.DataLoader(train_data, batch_size=10, shuffle=True, num_workers=2)
8
9 test_data = torchvision.datasets.CIFAR10(root='./cifar-10-python/test', train=False, download=True, transform=transform)
10 test_loader = torch.utils.data.DataLoader(test_data, batch_size=10, shuffle=False, num_workers=2)
```

Files already downloaded and verified  
Files already downloaded and verified

```
1 loss = nn.CrossEntropyLoss()
2 optimizer = optim.SGD(model.parameters(), lr=0.001, momentum=0.9)
3
4 epoch_num = 10
5
6 train_acc = [0] * epoch_num
7 train_error = [0] * epoch_num
8
9 print('Start')
10 for times in range(epoch_num):
11     correct_num = 0
12     total_num = 0
13     for i, data in enumerate(train_loader, 0):
14         # if random.randint(1,1000) % 1000 != 0: # remove
15         #     continue
16         if (i% 1000 == 0):
17             print(i)
18
19         inputs, labels = data
20
21         optimizer.zero_grad()
22         outputs = model(inputs)
23         current_loss = loss(outputs, labels)
24         current_loss.backward()
25         optimizer.step()
26         _, result = torch.max(outputs.data, 1)
27
28         correct_num = correct_num + (result == labels).sum().item()
29         total_num = total_num + labels.size(0)
30
31         train_error[times] = train_error[times] + current_loss.data.item()
32
33     train_acc[times] = correct_num/total_num
34     print(str(times) + ": acc is " + str(train_acc[times]) + ", error is " + str(train_error[times]/total_num))
35
36 print('Done')
```

```

1 train_error2 = train_error.copy()
2 for i in range(len(train_error)):
3     train_error[i] = train_error[i]/50000

```

```

1 # run on test data
2
3 loss = nn.CrossEntropyLoss()
4 optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
5
6 test_correct_num = 0
7 test_total_num = 0
8 test_error = 0
9 for i, data in enumerate(test_loader, 0):
10     # if random.randint(1,1000) % 1000 != 0: # remove
11     #     continue
12     if (i%1000 == 0):
13         print(i)
14
15     inputs, labels = data
16
17     optimizer.zero_grad()
18     outputs = model(inputs)
19     current_loss = loss(outputs, labels)
20     _, result = torch.max(outputs.data, 1)
21
22     test_correct_num = test_correct_num + (result == labels).sum().item()
23     test_total_num = test_total_num + labels.size(0)
24     test_error = test_error + current_loss.data.item()
25
26 print("Test data acc is " + str(test_correct_num/test_total_num) + ", error is " + str(test_error/test_total_num))

```

0  
Test data acc is 0.6791, error is 0.0927662309229374

```

1 plt.subplots(figsize=(7,5))
2 plt.plot(range(epoch_num), train_acc, label="train set")
3 plt.ylabel('accuracy')
4 plt.xlabel('epoch')
5 plt.legend()
6 plt.show()
7
8 plt.subplots(figsize=(7,5))
9 plt.plot(range(epoch_num), train_error, label="train set")
10 plt.ylabel('cross-entropy loss')
11 plt.xlabel('epoch')
12 plt.legend()
13 plt.show()

```

```

1 def print_filter():
2     print(model.layer1.weight.detach().shape)
3     layer = model.layer1.weight.detach().permute(0,2,3,1)
4     layer = layer - layer.min()
5     layer = layer / layer.max()
6     fig=plt.figure(figsize=(15,5))
7     for i in range(64):
8         f = layer[i]
9         fig.add_subplot(4,16, i+1)
10        plt.imshow(f)
11    plt.show()
12
13 print_filter()

```

torch.Size([64, 3, 11, 11])



## Part 2 batch normalization


### Structure

The new structure is defined as below, other code remains the same.

```
1  # input image: 3 * 32 * 32
2
3  class cnn(nn.Module):
4      def __init__(self):
5          super(cnn, self).__init__()
6          self.layer1 = nn.Conv2d(3, 64, 11)
7          self.pool1 = nn.MaxPool2d(2, 2)
8
9          self.layer2 = nn.Conv2d(64, 128, 3)
10         self.layer3 = nn.Conv2d(128, 128, 3)
11
12         self.pool2 = nn.AvgPool2d(3, 2)
13         self.fc1 = nn.Linear(128*3*3, 512)
14         self.fc2 = nn.Linear(512, 10)
15
16         self.batch1 = nn.BatchNorm2d(64)
17         self.batch2 = nn.BatchNorm2d(128)
18         self.batch3 = nn.BatchNorm2d(128)
19
20     def forward(self, x):
21         x = self.pool1(self.layer1(x))
22         x = self.batch1(x)
23         x = F.relu(self.layer2(x))
24         x = self.batch2(x)
25         x = F.relu(self.layer3(x))
26         x = self.batch3(x)
27         x = self.pool2(x)
28         x = x.reshape(-1, 128 * 3 * 3)
29         x = F.relu(self.fc1(x))
30         x = self.fc2(x)
31         return x
32
33 model = cnn()
```

## Compare training loss

```
1 plt.subplots(figsize=(7,5))
2 plt.plot(range(epoch_num), train_error, label="without batch")
3 plt.plot(range(epoch_num), train_error_batch, label="with batch")
4 plt.ylabel('cross-entropy loss')
5 plt.xlabel('epoch')
6 plt.legend()
7 plt.show()
```





## Final test error

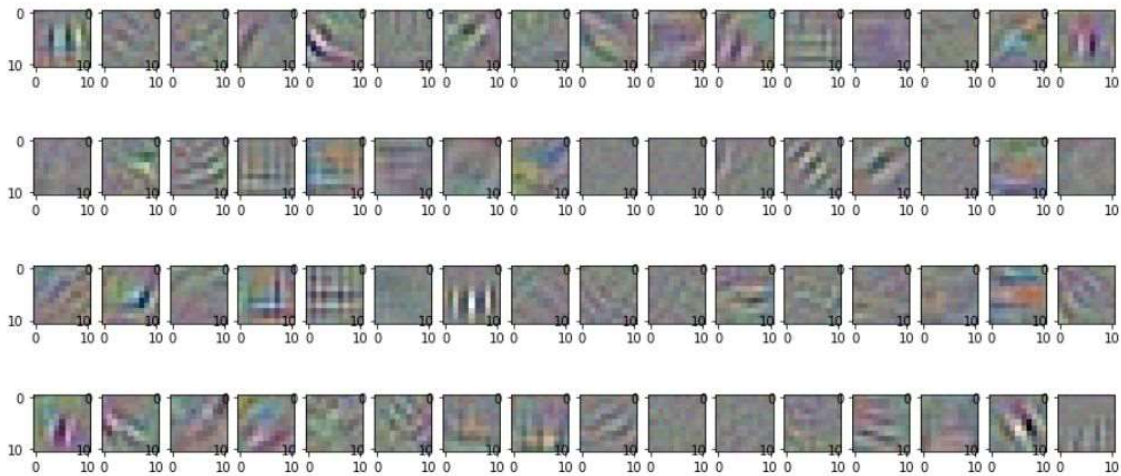
The final test error is 0.7016, which is better than the previous 0.6791.

```
1  # run on test data
2
3  loss = nn.CrossEntropyLoss()
4  optimizer = optim.SGD(model.parameters(), lr=0.01, momentum=0.9)
5
6  test_correct_num = 0
7  test_total_num = 0
8  test_error = 0
9  for i, data in enumerate(test_loader, 0):
10     # if random.randint(1,1000) % 1000 != 0: # remove
11     #     continue
12     # if (i % 1000 == 0):
13     #     print(i)
14
15     inputs, labels = data
16
17     optimizer.zero_grad()
18     outputs = model(inputs)
19     current_loss = loss(outputs, labels)
20     _, result = torch.max(outputs.data, 1)
21
22     test_correct_num = test_correct_num + (result == labels).sum().item()
23     test_total_num = test_total_num + labels.size(0)
24     test_error = test_error + current_loss.data.item()
25
26  print("Test data acc is " + str(test_correct_num/test_total_num))
```

Test data acc is 0.7016

## Filters

```
1 def print_filter():
2     # print(model.layer1.weight.detach().shape)
3     layer = model.layer1.weight.detach().permute(0,2,3,1)
4     layer = layer - layer.min()
5     layer = layer / layer.max()
6     fig=plt.figure(figsize=(15,5))
7     for i in range(64):
8         f = layer[i]
9         fig.add_subplot(4,16, i+1)
10        plt.imshow(f)
11    plt.show()
12
13 print_filter()
```




## Part 3 Optimize

By replacing the 11x11 filters with 3x3 filters and adding one more layer with convolution(128, 256) in the end, it will give a test result 0.7182 which is better than the previous result 0.7016 and 0.6791.

## Problem 4

```
1 for index in reversed(output2):
2     print(index_name[index], end = "")
3     print(" probability is ", end = "")
4     print(np.exp(output[0][index].detach().numpy())/prob)
5
6 imgplot = plt.imshow(img)
7 plt.show()
```

bell\_pepper probability is 0.9002557140463503



The rest is too hard :(

## Problem 5

ht is a three-dimension vector, two number represent two input unit and one number represent the carry on bit. Let first two number represent two input and last represent the carry on.

For current operation, we only need the carry on bit from the previous result, so

$$W = [[0,0,0] \ [0,0,0] \ [0,0,1]]$$

$$\text{We want } [0,0]*U = [0,0,0], [1,1]*U = [1,1,0], [1,0]*U = [1,0,0], [0,1]*U = [0,1,0]$$

$$\text{So } U = [[1,0], [0,1], [0,0]]$$

$$h_0 = [0,0,0] \text{ as there is no previous value to add.}$$

$$b_y = 0 \text{ as there is no bias}$$

$$b_h = [0,0,0]$$

$$v_t = [1,1,1], \text{ as we need to keep information from all three values.}$$